

Mostafa Rezaee

Ph.D. in Data Science

ML Engineer

Data Scientist

AI Scientist

GitHub: github.com/0-mostafa-rezaee-0

LinkedIn: linkedin.com/in/mostafa-rezaee/

Google Scholar: t.ly/JPyRL

Email: mostam@bgsu.edu

Mobile: 4193157481

Address: USA, Open to Relocation

HIGHLIGHTS

- **Innovative Data Solutions:** Spearheaded the development of the [SaveBirds.app](#), drastically reducing data preparation and analysis time by 99%, enabling ecologists to access comprehensive bird species data efficiently.
- **Leadership in Data Science:** Led a team of over 100 data analysts in a nationwide data digitization project, transforming the educational ecosystem for over a million students and significantly enhancing data accuracy and real-time updates.
- **Advanced Machine Learning Expertise:** Achieved significant improvements in predictive accuracy and processing efficiency in various projects, including a 26% improvement in gait speed prediction accuracy and a 45% reduction in data processing time using advanced ML techniques.
- **Publications & Research:** Authored 34 publications, including 12 journal articles and 22 full conference papers, showcasing expertise in computational studies. Led 7 projects, notably the [SaveBirds.app](#), collaborating with 10 researchers from 6 institutions.
- **GitHub:** Curated and maintained an extensive portfolio of repositories demonstrating technical expertise and open-source contributions across diverse domains, including LLMs, AI Engineering, ML Engineering, Recommender Systems, Causal Inference, A/B Testing, Physics & Engineering, and Personal Branding.

EXPERIENCE

Machine Learning Engineer

Stealth Startup, San Francisco, California

May 2024 - Present

- Deployed scalable ML models as REST APIs using FastAPI and Docker, integrating PostgreSQL for persistent prediction storage.
- Implemented real-time monitoring with Prometheus and Grafana, tracking API performance, model predictions, and system health.
- Automated ML workflows with CI/CD pipelines, enabling seamless model updates, containerized deployments, and continuous integration using GitHub Actions.

Data Scientist

SaveBirds, Bowling Green, Ohio

Sep 2019 - Apr 2024

- 99% reduction in data preparation and analysis time—cutting it from 90 days to just 1 minute—by developing the [SaveBirds.app](#), enabling ecologists and conservationists without coding skills to access 56 years of 800 bird species from 300,000 locations across North America.
- 99% reduction in atlas creation time—cutting it from 180 days to under 10 hours—by creating the Bird Atlas Generator (BAG), making it accessible without advanced GIS expertise.
- Automated key biodiversity metrics calculation enabling rapid, data-driven conservation decisions for 40,000 Protected Areas and supporting the [\\$75 billion](#) wildlife-watching industry.
- Supported 7 projects involving 10 researchers from 6 institutions.

AI Scientist

Sanofi, Cambridge, Massachusetts

Jun 2022 - Aug 2022

- **Boosted Gait Speed Accuracy:** Improved prediction accuracy from 70% to 88% (a 26% relative improvement) using LSTM networks on accelerometer data.
- **Refined Step Segmentation:** Improved classification accuracy from 86% to 94% reducing error rates from 14% to 6% (a 57% relative improvement) through Random Forest, SVM, and advanced signal processing techniques.

- **Reduced Drift in Step Length Calculation:** Enhanced stride length estimation accuracy from 75% to 97% (a 29% relative improvement) and reduced measurement error from 6.5 cm to 5.1 cm using advanced Kalman filtering techniques.
- **Minimized False Positives in Sway Detection:** Improved accuracy in detecting sway from 65% to 95% (a 46% relative improvement) and reduced false alarms from 22% to 15% using Bayesian filtering, adaptive thresholding, and time-series anomaly detection algorithms.
- **Optimized Data Pipeline:** Accelerated preprocessing of 1.2 million accelerometer data points, reducing processing time from 9.2s to 5.1s (a 45% relative improvement) by implementing Apache Spark and advanced feature engineering techniques.

Lead Data Scientist & Deputy Director of Research and Technology

Farabi Institute, Ahvaz, Khuzestan

Sep 2013 - Aug 2019

- Led the provincial implementation of a nationwide data digitization project, transforming the educational ecosystem for over 1,000,000 students, 40,000 classrooms, and 76,000 teachers across 40 districts.
- Supervised a team of 100+ data analysts directly reporting to me, ensuring standardized, high-quality data collection and analysis at scale.
- Enhanced data accuracy by 40% and established real-time updates, allowing instant visibility into changes in student, teacher, and school profiles.
- Applied advanced regression, classification, and time series analyses to derive actionable insights, guiding data-driven policy decisions for senior authorities.
- Developed an automated alert system that identified significant performance shifts, prompting timely interventions and continuous improvement throughout the education system.

Adjunct Professor of Machine Learning and Computational Studies

Payame Noor University (PNU), Ahvaz, Khuzestan

Sep 2013 - Aug 2019

- Empowered 600+ students over six years by teaching a 3-credit Computer Programming course each semester, covering Python and Machine Learning to 50+ senior and master's students per class.
- Supervised 100+ senior students' final projects, designing specialized Machine Learning projects that laid the foundation for their master's research and careers in AI.
- Shaped the future of AI professionals by inspiring students to integrate Machine Learning into their master's projects, fostering a new generation of ML practitioners and researchers.

EDUCATION

Data Science, Ph.D. | *Bowling Green State University, Ohio.*

2024

SKILLS

Demonstrates strong analytical, creative, and problem-solving skills. Proficient in communicating complex ideas to diverse audiences. Tech-savvy with a keen ability to stay updated with the latest data programming tools and methodologies.

With a PhD in Data Science, focused on applied mathematics, I merge rigorous theory with production pragmatism. My day-to-day work ranges from deriving bespoke algorithms grounded in calculus and linear algebra to architecting full-stack data analytics pipelines. I excel at integrating heterogeneous data sources, orchestrating ETL processes, and automating model-deployment integration with Docker- and Kubernetes-based CI/CD. Whether I am building probabilistic models, fine-tuning transformer LLMs, or optimizing SQL-based querying strategies, I always select the right statistical and visualization tools—from PyTorch and MLflow to Tableau and Power BI—to deliver insights that drive revenue, not just reports.

Data Science A/B Testing, Causal Inference, Recommender Systems

LLMs & Generative AI Transformers, RAG, LangChain, OpenAI, Claude, Prompt Engineering, Ray

AI Engineering FastAPI, Streamlit, Gradio, HuggingFace, ONNX Runtime, Triton, W&B

ML Engineering Docker, Kubernetes, MLflow, Kubeflow, Apache Airflow, SageMaker

Deep Learning Computer Vision, Natural Language Processing (NLP), HuggingFace, PyTorch, TensorFlow

Machine Learning scikit-learn, caret, XGBoost, LightGBM, Random Forest

Programming Languages Python, SQL, R, C++, FORTRAN

Statistics Hypothesis Testing, Bayesian Methods, Experimental Design, Probability Theory

Database PostgreSQL, MySQL, MongoDB, Redis, Elasticsearch

DevOps/MLOps CI/CD Pipelines, GitHub Actions, Jenkins, Monitoring (Prometheus, Grafana)

Cloud Platforms AWS, Azure, Google Cloud

Time Series Analysis Transformers, TCN, Prophet, LSTM, Statsmodels, ARIMA, SARIMA, TSA

Big Data Apache Spark, Hadoop, Apache Hive, Presto, Apache Flink, Dask

Software Engineering Git, Design Patterns, Test-Driven Development, PyTest, Unit Testing

Visualization Matplotlib, Seaborn, Plotly, Tableau, PowerBI